

Exercise 1 (Linear Regression I)

We consider the `agefat` dataset from library `HSAUR2` and want to investigate the variable body fat (`fat`). The data can be loaded with:

```
# install.packages("HSAUR2")
library(HSAUR2)

## Loading required package: tools

data("agefat")
dat = agefat
head(dat)


##   age  fat gender
## 1  24 15.5  male
## 2  37 20.9  male
## 3  41 18.6  male
## 4  60 28.0  male
## 5  31 34.7 female
## 6  39 30.2 female
```

- (a) Investigate the relationship between age (`age`) and body fat percentage (`fat`) as well as between gender (`gender`) and body fat percentage graphically.

```
par(mfrow=c(1,2))

# Since age and body fat are continuous, we use a scatterplot
plot(dat$age, dat$fat,
      ylab="Body fat",
      xlab="Age")

# Since gender is dichotomous, we use a boxplot
boxplot(dat$fat~dat$gender,
        names = c('female','male'),
        ylab='Body fat',
        xlab='Gender')
```



- (b) Fit a linear regression model using the `lm()` function. Use age and gender as covariates. Interpret the estimates for the intercept, age and gender. (**R-Hint:** to fit the model use `mod <- lm()`. To consider the results use `summary(mod)`)

```
mod = lm(fat ~ age + gender, data=dat)
summary(mod)

##
## Call:
## lm(formula = fat ~ age + gender, data = dat)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -8.4691 -2.0505  0.0276  2.1442  8.2773
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  19.64786    4.10781   4.783 8.92e-05 ***
## age          0.26556    0.07953   3.339  0.00297 **
## gendermale  -10.54892    2.09140  -5.044 4.74e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.644 on 22 degrees of freedom
## Multiple R-squared:  0.7375, Adjusted R-squared:  0.7137
## F-statistic: 30.91 on 2 and 22 DF, p-value: 4.069e-07
```

```
# Intercept: mean body fat percentage if both explanatory
# variables are 0, i.e. if age=0 and gender=female (response).

# Age: If age increases by one unit, the mean body fat percentage increases
# by 0.2 units given gender stays constant. That is, with increasing age,
# body fat percentage increases if gender is the same.

# Gender: The mean body fat percentage is about 10 units smaller for males
# than for females, given that age doesn't change.
```


- (c) Check the model assumptions using a Tukey-Ascombe and a normal QQ-plot (**R-Hint:** To get the fitted values and the residuals for the Tukey-Ascombe plot, you can use `fitted(mod)` and `resid(mod)`. For the QQ-plot use the function `qqPlot()` from library `car`).

```
par(mfrow=c(1,2))

# TA plot: fitted vs. residuals
plot(fitted(mod), resid(mod),
     main='Tukey-Ascombe Plot',
     ylab='Residuals',
     xlab='Fitted')
abline(a=0, b=0, col='blue')

# Normal QQ-plot with the residuals
library(car)

## Loading required package: carData
qqPlot(resid(mod), dist = "norm",
       mean = mean(resid(mod)),
       sd = sd(resid(mod)),
       ylab='Residual quantiles',
       xlab='Normal quantiles',
       main="QQ-Plot")
```



```
## [1] 10 18
# The model assumptions hold.
```

- (d) Write down the equation of the model ($Y = \dots$). Then predict the mean body fat percentage for a 40 year old woman. You can calculate it by hand or by using the R-function `predict()`.

```
# The formular is:
#  $Y = 19.65 + 0.27 \cdot \text{age} - 10.55 \cdot \text{gender}$ 

# The mean body fat percentage for a 40 year old woman is:
#  $Y = 19.65 + 0.27 \cdot 40 - 10.55 \cdot 0 = 30.45$ 
predict(mod, newdata=data.frame(age=40, gender="female"))

##          1
## 30.27037
```

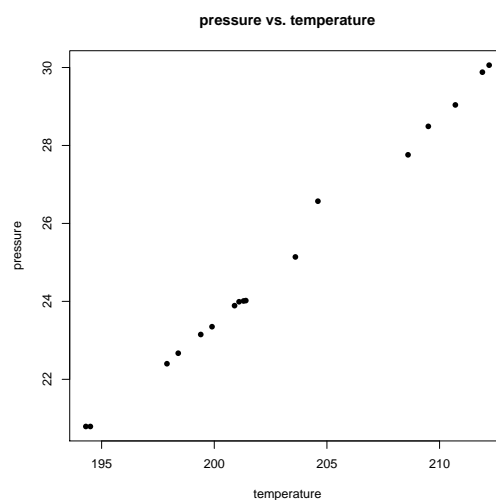
Exercise 2 (Linear regression II)

The data set of Forbes lists the boiling point of water (in $^{\circ}\text{F}$) and the atmospheric pressure (in inches of mercury) at different places in the alps. We want to investigate the association between the temperature (Temp) and the pressure (Press). You can read the data into R using:

```
url = "https://polybox.ethz.ch/index.php/s/uZJZavllfYbxldy/download"
dat = read.table(url, header = TRUE)
```

- (a) Investigate the relationship between pressure and temperature graphically. Is it reasonable to fit a linear regression model?

```
plot(dat$Temp, dat$Press, pch=16,
     main='pressure vs. temperature',
     ylab='pressure',
     xlab='temperature')
```



```
# If we consider the data it seems to be reasonable to fit a linear regression
# model. We see, that the atmospheric pressure increases with the increase
# of the boiling point of water.
```

- (b) Perform a linear regression. Investigate the influence of the temperature (covariate) on the pressure (outcome).

```
mod = lm(Press~Temp, data = dat)
summary(mod)

##
## Call:
## lm(formula = Press ~ Temp, data = dat)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.25717 -0.11246 -0.05102  0.14283  0.64994
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -81.06373     2.05182  -39.51  <2e-16 ***
## Temp         0.52289     0.01011   51.74  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2328 on 15 degrees of freedom
## Multiple R-squared:  0.9944, Adjusted R-squared:  0.9941
## F-statistic: 2677 on 1 and 15 DF,  p-value: < 2.2e-16

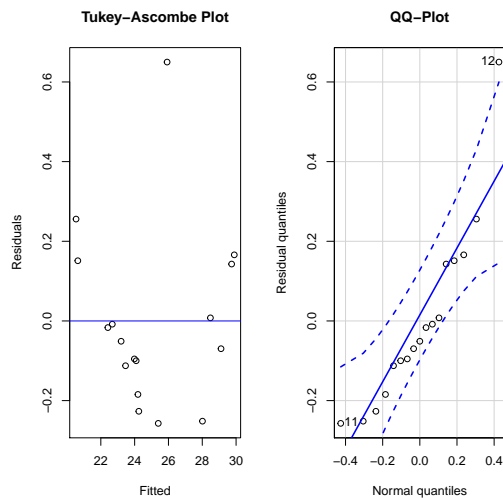
# If the temperature increases by one point, the atmospheric pressure
# increases by 0.5 inches. That is, the higher the temperature for the
# boiling point of the water, the higher the atmospheric pressure.
```

- (c) Generate a Tukey-Anscombe plot and a normal Q-Q plot of the residuals. Are there any hints that the model assumptions are violated?

```
par(mfrow=c(1,2))

# TA plot: fitted vs. residuals
plot(fitted(mod), resid(mod),
     main='Tukey-Ascombe Plot',
     ylab='Residuals',
     xlab='Fitted')
abline(a=0, b=0, col='blue')

# Normal QQ-plot with the residuals
library(car)
qqPlot(resid(mod), dist = "norm",
       mean = mean(resid(mod)),
       sd = sd(resid(mod)),
       ylab='Residual quantiles',
       xlab='Normal quantiles',
       main="QQ-Plot")
```



```
## [1] 12 11

# The Tukey Ascombe plot shows that there are some violations of the
# assumptions. The residuals are not in a band around the horizontal
# line.
# The QQ-plot looks good but it shows an outlier (observation 12)
# which we should investigate further.
```

- (d) To get a better fit, we transform the outcome variable Press. Calculate the log of the variable and fit a linear regression model with the transformed outcome variable. Calculate a Tukey-Ascombe and a QQ-plot. What about the model assumptions?


```
# calculate the log
dat$log_Press = log(dat$Press)
log_mod = lm(log_Press ~ Temp, data = dat)

par(mfrow=c(1,2))

# TA plot: fitted vs. residuals
plot(fitted(log_mod), resid(log_mod),
     main='Tukey-Ascombe Plot',
     ylab='Residuals',
     xlab='Fitted')
abline(a=0, b=0, col='blue')

# Normal QQ-plot with the residuals
```

```
library(car)
qqPlot(resid(log_mod), dist = "norm",
       mean = mean(resid(log_mod)),
       sd = sd(resid(log_mod)),
       ylab='Residual quantiles',
       xlab='Normal quantiles',
       main="QQ-Plot")
```



```
## [1] 12 14
```

```
# The Tukey Ascombe plot shows no violation anymore.
# However, the outlier is still a problem.
```

- (e) Identify and remove the outlier. Calculate a Tukey-Ascombe and a QQ-plot. What about the model assumptions now?

```
# We saw that the observation 12 is the outlier:
dat_new = dat[-12,]

mod_new = lm(log_Press ~ Temp, data = dat_new)

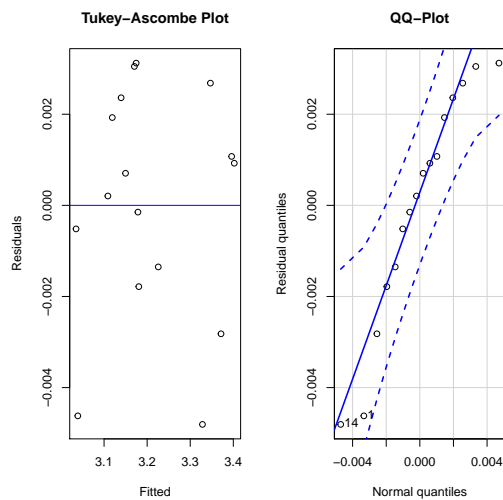
par(mfrow=c(1,2))

# TA plot: fitted vs. residuals
plot(fitted(mod_new), resid(mod_new),
```



```
main='Tukey-Ascombe Plot',
ylab='Residuals',
xlab='Fitted')
abline(a=0, b=0, col='blue')

# Normal QQ-plot with the residuals
library(car)
qqPlot(resid(mod_new), dist = "norm",
mean = mean(resid(mod_new)),
sd = sd(resid(mod_new)),
ylab='Residual quantiles',
xlab='Normal quantiles',
main="QQ-Plot")
```



```
## 14 1
## 13 1

# All assumptions hold now.
```