

Exercise 1

We consider the dataset diabetes (Efron, Hastie, Johnstone and Tibshirani (2003) "Least Angle Regression" Annals of Statistics) from the package lars. Ten baseline variables age, sex, body mass index (bmi), average blood pressure (map) and six blood serum measurements (tc, ldl, hdl, tch, ltg, glu), as well as disease progression one year after baseline (y), were obtained for n=442 diabetes patients. The baseline data is stored in x while a model matrix including interactions between baseline measurements is stored in x2. Here, we aim to predict the disease progression, one year after baseline based on the matrix x2. You can access the data via

```
# install.packages("lars")
library(lars)

## Warning: package 'lars' was built under R version 3.4.4
## Loaded lars 1.2

data("diabetes")
```

- (a) Split the data set into a training and a test set. Sample 70% of the data to the training set, 30% to the test set. Set the seed to 100 (set.seed(100)).

```
set.seed(100)

# get index for 70% of the data
train_idx = sample(1:nrow(diabetes), 0.7*nrow(diabetes))
train <- diabetes[train_idx, c("y", "x2")]
test <- diabetes[-train_idx, c("y", "x2")]
```

- (b) Fit a linear regression model based on the training data and check the model assumptions. Is it important that all the assumptions are met? Now, use the model for prediction on the test data. Calculate the test error in terms of the mean squared error (MSE) and the mean absolute percentage error (MAPE) using OLS. Consider the predicted vs. the observed values on the test data.

```
# linear regression
mod <- lm(y~x2, data=train)
summary(mod)

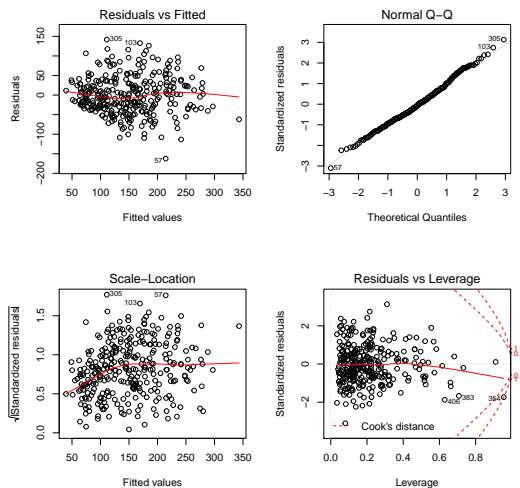
##
## Call:
```

```
## lm(formula = y ~ x2, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -162.396  -31.078   -3.148   29.643  141.968
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    151.285      3.211  47.109 < 2e-16 ***
## x2age           12.901      85.143   0.152  0.87969
## x2sex          -190.338      82.425  -2.309  0.02177 *
## x2bmi           377.872     115.199   3.280  0.00119 **
## x2map           231.733     100.174   2.313  0.02154 *
## x2tc          -19648.891  64386.998  -0.305  0.76050
## x2ldl          17211.983  56584.458   0.304  0.76125
## x2hdl           7061.042  24070.722   0.293  0.76951
## x2tch           -16.251     345.102  -0.047  0.96248
## x2ltg           7243.646  21169.619   0.342  0.73252
## x2glu            35.088      90.743   0.387  0.69933
## x2age^2          99.230      86.208   1.151  0.25084
## x2bmi^2          98.643     122.428   0.806  0.42119
## x2map^2          28.890      90.247   0.320  0.74915
## x2tc^2          136.230   9376.141   0.015  0.98842
## x2ldl^2         -971.065   6818.355  -0.142  0.88687
## x2hdl^2        -847.828   2429.442  -0.349  0.72740
## x2tch^2         1310.289    780.017   1.680  0.09427 .
## x2ltg^2         1158.445   1846.271   0.627  0.53095
## x2glu^2          90.027    110.335   0.816  0.41533
## x2age:sex        156.478     94.522   1.655  0.09912 .
## x2age:bmi       -90.200    107.401  -0.840  0.40182
## x2age:map        38.504     98.186   0.392  0.69529
## x2age:tc       -712.407    845.810  -0.842  0.40046
## x2age:ldl       338.366    677.871   0.499  0.61812
## x2age:hdl       439.196    383.092   1.146  0.25273
## x2age:tch       308.114    283.590   1.086  0.27834
## x2age:ltg       329.231    298.426   1.103  0.27102
## x2age:glu        61.207    100.420   0.610  0.54275
## x2sex:bmi        49.579    109.161   0.454  0.65010
## x2sex:map        76.282    100.337   0.760  0.44783
## x2sex:tc       1732.571    959.500   1.806  0.07220 .
```

```
## x2sex:ldl      -1461.008      768.291      -1.902      0.05840 .
## x2sex:hdl      -615.455      424.122      -1.451      0.14803
## x2sex:tch      -192.215      258.983      -0.742      0.45869
## x2sex:ltg      -527.721      335.377      -1.574      0.11689
## x2sex:glu       88.975       91.273       0.975      0.33061
## x2bmi:map       76.129      124.727       0.610      0.54219
## x2bmi:tc       563.706      997.144       0.565      0.57238
## x2bmi:ldl     -320.287      848.540      -0.377      0.70616
## x2bmi:hdl     -463.675      492.374      -0.942      0.34727
## x2bmi:tch     -366.729      348.217      -1.053      0.29331
## x2bmi:ltg     -158.937      374.380      -0.425      0.67155
## x2bmi:glu     -18.519      131.696      -0.141      0.88828
## x2map:tc      -692.624     1176.219      -0.589      0.55650
## x2map:ldl      660.140     1001.921       0.659      0.51060
## x2map:hdl      231.736      503.491       0.460      0.64574
## x2map:tch     -39.519      280.077      -0.141      0.88791
## x2map:ltg      257.768      423.731       0.608      0.54354
## x2map:glu     -152.860      114.685     -1.333      0.18382
## x2tc:ldl      1253.551     15388.076       0.081      0.93514
## x2tc:hdl      1649.567      5575.221       0.296      0.76758
## x2tc:tch     -1198.123      2457.628      -0.488      0.62633
## x2tc:ltg      2561.147     14638.499       0.175      0.86126
## x2tc:glu       81.902      919.932       0.089      0.92913
## x2ldl:hdl     -2201.100      4634.229      -0.475      0.63524
## x2ldl:tch       51.932      2093.967       0.025      0.98023
## x2ldl:ltg     -2395.625     12116.436      -0.198      0.84343
## x2ldl:glu     -240.020       798.622      -0.301      0.76402
## x2hdl:tch      1257.690      1301.771       0.966      0.33493
## x2hdl:ltg     -1503.305      5243.996      -0.287      0.77461
## x2hdl:glu      231.672      417.732       0.555      0.57968
## x2tch:ltg     -51.091      840.475      -0.061      0.95158
## x2tch:glu      486.218      299.345       1.624      0.10561
## x2ltg:glu       52.390      380.447       0.138      0.89059
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 54.58 on 244 degrees of freedom
## Multiple R-squared:  0.5908, Adjusted R-squared:  0.4835
## F-statistic: 5.506 on 64 and 244 DF,  p-value: < 2.2e-16

# check the model assumptions
```

```
par(mfrow=c(2,2))  
plot(mod)
```



```
# The model assumptions are not violated. However, even if  
# they were, it wouldn't be so important because we aim to  
# use the model for prediction. If it predicts the test data  
# well we don't care about the violations on the training data
```

```
# predict the data on the test set  
predLM <- predict(mod, newdata = test)
```

```
# MSE and MAPE for the test data  
mean((test$y - predLM)^2)
```


```
## [1] 3896.808
```

```
mean(abs((predLM - test$y))/test$y)
```

```
## [1] 0.3926821
```

```
# plot the observed vs the fitted
```

```
par(mfrow=c(1,1))  
plot(predLM, test$y, main="Predicted vs observed",  
      xlab="predicted", ylab="observed")  
abline(0,1, col='green', lty=2, lwd=2)
```



- (c) Now, we aim to predict the test data using ridge regression. Recall what ridge regression is doing and what's the impact of λ . Perform a cross validation to find the best parameter λ based on the training data. Is the model, fitted with the optimal parameter λ , a better prediction model for the test data compared to the linear regression model? Plot the predicted vs. the observed values.

```
library(glmnet)

## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-16
set.seed(100)

# fit a ridge regression
mod_ridge = cv.glmnet(x=train$x2, y=train$y, alpha=0)

# In ridge regression we extend the optimization
# objective using a penalty term for large coefficients,
# given by the sum of squared coefficients. Lambda is a
# tuning parameter that determines the contribution of the
# penalty term to the equation. The penalty leads to
# coefficients, shrunk towards zero. These coefficients
# are less optimal on the training data but they lead
# to a better prediction performance on new test data.
```

```
# best
lambda_ridge <- mod_ridge$lambda.min
lambda_ridge

## [1] 30.85116

# predict the data
predRidge <- predict(mod_ridge, newx = test$x2, s = lambda_ridge)

# MSE and MAPE on the test data
mean((test$y - predRidge)^2)


## [1] 2904.235

mean(abs((predRidge - test$y))/test$y)

## [1] 0.3840076

# The MSE and the MAPE are smaller compared to the MSE and
# the MAPE of the linear regression indicating better predictions
# on the test data.

# plot the observed vs the fitted
par(mfrow=c(1,1))
plot(predRidge, test$y, main="Predicted vs observed",
      xlab="predicted", ylab="observed")
abline(0,1, col='green', lty=2, lwd=2)
```




```
# The slope of the predicted data seems to be close to one  
# indicating a good prediction on the test data
```

- (d) Now, we aim to predict the test data using a lasso regression. What's the difference to the ridge regression? Perform a cross validation on the training data to find the best parameter λ (`cv.glmnet(..., alpha=1)`). Is the model, fitted with the optimal parameter λ , a better prediction model than the linear and the ridge regression? Plot the predicted vs the observed values.

```
set.seed(100)  
  
# fit a lasso regression  
mod_lasso <- cv.glmnet(x=train$x2, y=train$y, alpha=1)  
  
# In Lasso and ridge regression we extend the optimization  
# objective using a penalty term for large coefficients.  
# Opposed to ridge regression, the penalty term is not the  
# L2 but the L1 norm of the coefficients. This leads to  
# shrunked coefficient estimates. These estimates are,  
# in contrast to the ridge regression estimates, more often  
# exactly equal to zero meaning, that the respective  
# predictor is removed from the model.  
  
# predict the best lambda  
lambda_lasso <- mod_lasso$lambda.min  
lambda_lasso  
  
## [1] 2.944893  
  
# predict the data in the test set  
predLasso <- predict(mod_lasso, newx = test$x2, s=lambda_lasso)  
  
# MSE and MAPE  
mean((test$y - predLasso)^2)  
  
## [1] 2891.84  
  
mean(abs((predLasso - test$y))/test$y)  
  
## [1] 0.3764373  
  
# The MSE and the MAPE are smaller than the ones from the  
# linear respectively ridge regression, indicating better
```

```
# predictions

# plot the observed vs the fitted
par(mfrow=c(1,1))
plot(predLasso, test$y, main="Predicted vs observed",
      xlab="predicted", ylab="observed")
abline(0,1, col='green', lty=2, lwd=2)
```



- (e) Calculate the predictions on the training data for each of the three models. Which model fits best? Do the results make sense?

```
# predict the results on the training data
predLM_train <- predict(mod, train)
predRidge_train <- predict(mod_ridge, newx = train$x2, s = lambda_ridge)
predLasso_train <- predict(mod_lasso, newx = train$x2, s = lambda_lasso)

# get the MSE
mean((train$y - predLM_train)^2)

## [1] 2352.198

mean((train$y - predRidge_train)^2)

## [1] 2716.099

mean((train$y - predLasso_train)^2)

## [1] 2791.764
```



```
# get the MAPE
mean(abs((predLM_train - train$y))/train$y)

## [1] 0.3385566

mean(abs((predRidge_train - train$y))/train$y)

## [1] 0.3815519

mean(abs((predLasso_train - train$y))/train$y)

## [1] 0.3842834

# Based on the training data, the linear model is better than
# the ridge regression which fits slightly better than the
# lasso regression.
# This makes sense because the least square estimates lead
# to the best and unbiased model w.r.t to the data used for
# fitting the model. By extending the optimization objective with
# the penalty term for large coefficients, we get shrinked
# coefficient estimates which are less optimal on the training data.
# However, they lead to better prediction performance on new test
# data.
```