

Biostatistics: Exercise 11

Beate Sick, Lisa Herzog

24.11.2020

Exercise 1: Random Forest for classification

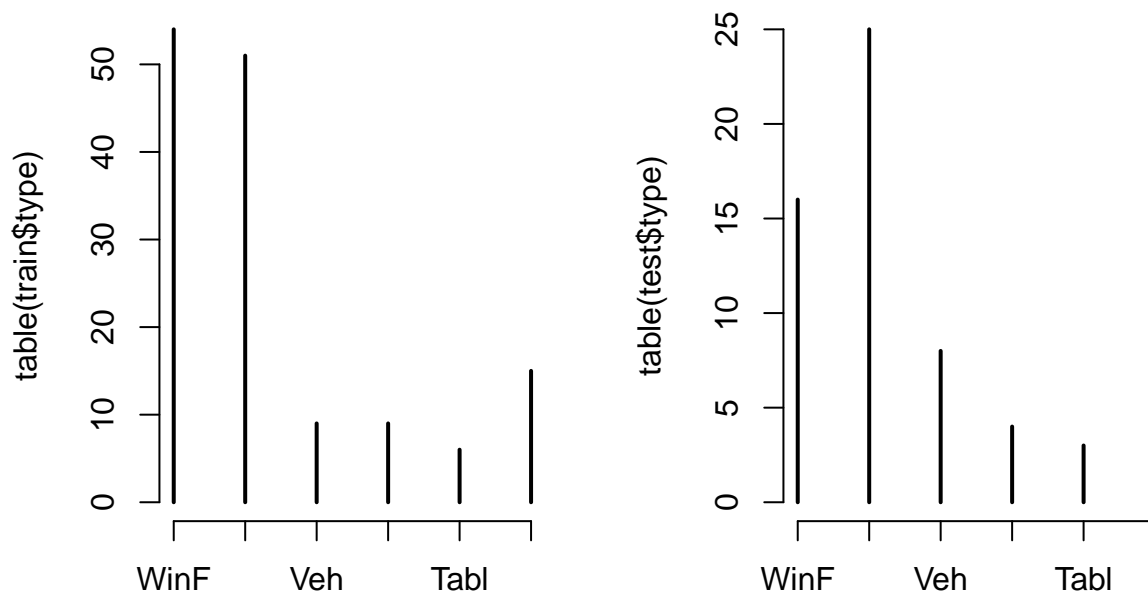
The goal in this exercise is to use a Random Forest for classification. The data set summarizes the chemical concentration of 9 different elements (e.g. Na and Mg) and each observation corresponds to one out of 6 classes corresponding to different types of glass fragments. You can download the training `train.fgl.RData` and the test data set `test.fgl.RData` from the Website.

- Load the training and the test data into R using the function `load()` and become acquainted with the data. Perform a descriptive analysis. Assess how the target variable `type` is distributed in the training and the test data and evaluate the pair-wise relationships between the explanatory variables and the target `type`. Comment on your analysis results.

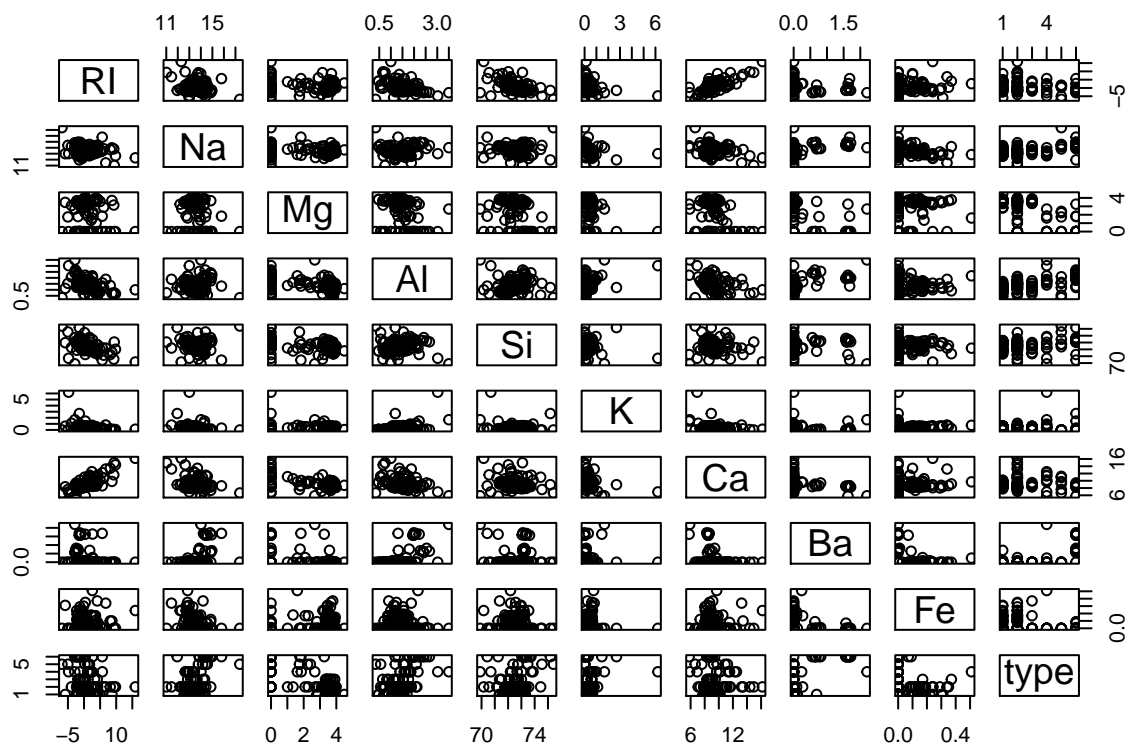
```
load(paste0(dir, "data/train.fgl.RData"))
load(paste0(dir, "data/test.fgl.RData"))

str(train)
## 'data.frame':   144 obs. of  10 variables:
##  $ RI   : num  3.01 -1.82 -0.58 -0.57 -0.44 ...
##  $ Na   : num  13.6 13.5 13.3 13.3 13.2 ...
##  $ Mg   : num  4.49 3.55 3.62 3.6 3.61 3.46 3.66 3.56 3.59 3.54 ...
##  $ Al   : num  1.1 1.54 1.24 1.14 1.05 1.56 1.27 1.27 1.31 1.23 ...
##  $ Si   : num  71.8 73 73.1 73.1 73.2 ...
##  $ K    : num  0.06 0.39 0.55 0.58 0.57 0.67 0.6 0.54 0.58 0.58 ...
##  $ Ca   : num  8.75 7.78 8.07 8.17 8.24 8.09 8.56 8.38 8.5 8.39 ...
##  $ Ba   : num  0 0 0 0 0 0 0 0 0 0 ...
##  $ Fe   : num  0 0 0 0 0 0.24 0 0.17 0 0 ...
##  $ type: Factor w/ 6 levels "WinF","WinNF",...: 1 1 1 1 1 1 1 1 1 1 ...

# visualize distribution of target variable
par(mfrow = c(1,2))
plot(table(train$type))
plot(table(test$type))
```



```
# from the barplots we see that the levels of the target variable are unbalanced.  
# The distribution of the target variable is similar in the train and test.  
  
# Get a visual impression of pair-wise relationship between variables and target  
pairs(train)
```



*# from the last column in the pairs plot we see that none of the explanatory
variables alone can separate the classes reasonably well.*

- Use the training data to fit a classification RF. Set the arguments to `importance=TRUE` for a later assessment of the importance of the different explanatory variables on the target and `ntree=1000`.
 - How large is the out-of-bag error over all classes?
 - Which class(es) are especially hard to classify correctly?
 - Which class(es) are most easy to classify correctly?

```
library(randomForest)
## Warning: package 'randomForest' was built under R version 4.0.3
## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
set.seed(0815)
rf1 = randomForest(type ~ ., data = train, ntree = 1000, importance = TRUE )
print(rf1)
##
## Call:
## randomForest(formula = type ~ ., data = train, ntree = 1000,      importance = TRUE)
##
##           Type of random forest: classification
##           Number of trees: 1000
## No. of variables tried at each split: 3
##
##           OOB estimate of error rate: 24.31%
```

```
## Confusion matrix:
##      WinF WinNF Veh Con Tabl Head class.error
## WinF    48    5  1  0  0  0  0.1111111
## WinNF   10   39  0  1  1  0  0.2352941
## Veh      8    1  0  0  0  0  1.0000000
## Con      0    3  0  5  0  1  0.4444444
## Tabl     1    1  0  0  4  0  0.3333333
## Head     0    2  0  0  0 13  0.1333333

# OOB estimate of the error rate is ~24%
# VEH is the hardest class to predict with an error of 100%
# WinF and Head appear to be easy to predict with low error of about 10%
```

- Use the trained RF to predict the classes in the test data. Determine the test confusion matrix, the accuracy and the misclassification rate. Comment on your results.

```
# predict the test data
rf1.pred = predict(rf1, newdata = test, type = "class")

# Confusion matrix
(table2 = table(rf1.pred, test$type))
##
## rf1.pred WinF WinNF Veh Con Tabl Head
## WinF    12    3  5  0  0  1
## WinNF     4   22  2  1  0  2
## Veh       0    0  1  0  0  0
## Con       0    0  0  3  0  0
## Tabl      0    0  0  0  3  0
## Head      0    0  0  0  0 11

# accuracy
accuracy2 = sum(diag(table2))/sum(table2)
accuracy2
## [1] 0.7428571


# The accuracy is ~74%.
# This corresponds to a misclassification rate of 26%: 1 - 0.74 = 0.26

# This is similar to the OOB misclassification rate in the training data.
# The OOB error in the training data is achieved based on trees in the RF that
# did not see the observations in the bootstrap-train-sample.
# Since the training and the test data come from the same distribution, we
# have expected similar misclassification rates in the training and test data set.
```

- Which explanatory variables are most important for the classification?

```
varImpPlot(rf1)
```

rf1



the concentration of Mg is most useful for the classification

Exercise 2: Random Forest versus lm for a regression model with continuous outcome

- The data set Boston is available in the package MASS. Load it and explore the help page to grab a minimal understanding of the data.

```
library(MASS)

data(Boston)
#help("Boston")

dim(Boston)
## [1] 506 14
# 506 obs., 14 variables

str(Boston)
## 'data.frame': 506 obs. of 14 variables:
## $ crim : num 0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn : num 18 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus : num 2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 ...
## $ chas : int 0 0 0 0 0 0 0 0 0 ...
## $ nox : num 0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 ...
## $ rm : num 6.58 6.42 7.18 7 7.15 ...
## $ age : num 65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis : num 4.09 4.97 4.97 6.06 6.06 ...
```

```
## $ rad      : int  1 2 2 3 3 3 5 5 5 5 ...
## $ tax      : num 296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 18.7 15.2 15.2 15.2 ...
## $ black    : num 397 397 393 395 397 ...
## $ lstat    : num 4.98 9.14 4.03 2.94 5.33 ...
## $ medv     : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...

# We see chas is not correctly coded:
Boston$chas <- as.factor(Boston$chas)
```

- Randomly split the data into two subsets, a training and a test data set, using the proportion of 70% - 30%.

```
# We randomly sample the row indexes of the data set:
set.seed(1)
idx.tr <- sample(x = 1:nrow(Boston), size = as.integer(0.7 * nrow(Boston)))

train <- Boston[idx.tr, ]
test  <- Boston[-idx.tr, ]
```

- Fit a regression model (once with `lm` and once with `randomForest`) with `medv` as target variable and all other variables as predictors. Fit the models using the training set.

```
# Linear Regression Model with "medv" as target and 13 covariates:
fit.lm <- lm(medv ~ ., data = train)
summary(fit.lm)
##
## Call:
## lm(formula = medv ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.1732  -2.7272  -0.5408   1.6594  23.9450
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  28.209249   6.094508   4.629 5.24e-06 ***
## crim         -0.070644   0.042379  -1.667 0.096439 .
## zn           0.036046   0.016006   2.252 0.024959 *
## indus        0.030055   0.069986   0.429 0.667874
## chas1        3.419418   0.933343   3.664 0.000288 ***
## nox         -14.582615   4.374454  -3.334 0.000952 ***
## rm           4.856324   0.493995   9.831 < 2e-16 ***
## age         -0.014469   0.015251  -0.949 0.343431
## dis         -1.411125   0.231606  -6.093 3.01e-09 ***
## rad          0.303706   0.075162   4.041 6.59e-05 ***
## tax         -0.013742   0.004219  -3.257 0.001238 **
## ptratio     -0.941851   0.154394  -6.100 2.88e-09 ***
## black        0.010698   0.003084   3.469 0.000590 ***
## lstat       -0.480384   0.061173  -7.853 5.37e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.593 on 340 degrees of freedom
## Multiple R-squared:  0.7778, Adjusted R-squared:  0.7693
```

```
## F-statistic: 91.52 on 13 and 340 DF,  p-value: < 2.2e-16

# we assume that the model assumptions for lm are fulfilled

# Now the RF for the continuous outcome medv:
library(randomForest)
fit.rf <- randomForest(medv ~ ., data = test, importance = TRUE)
fit.rf
##
## Call:
## randomForest(formula = medv ~ ., data = test, importance = TRUE)
##              Type of random forest: regression
##              Number of trees: 500
## No. of variables tried at each split: 4
##
##              Mean of squared residuals: 21.27425
##              % Var explained: 68.43



# for the RF there are no model assumptions that need to be checked
```

- Get the predictions for the test data using the fitted models (lm and rf) and plot the observed medv values in the test set versus the predicted values. Based on the plot – how do both models compare?

```
par(mfrow = c(1, 2))

# For the lm Model:
preds.lm <- predict(fit.lm, test)
plot(preds.lm, test$medv)
abline(0, 1, col = "red")

#For the rf:
preds.rf <- predict(fit.rf, test)
plot(preds.rf, test$medv)
abline(0, 1, col = "red")
```



*# All models yield quite unbiased predictions
but the variance of the RF predictions is much lower*

- Calculate the mean squared error (MSE) of these predictions on the test set. Is the MSE better for lm or for the random forest? (Hint: $MSE = \frac{1}{n} \sum_{i=1}^n (\hat{Y}_i - Y_i)^2$)

```
(mse.lm <- mean((preds.lm - test$medv)^2))  
## [1] 27.31196
```

```
(mse.rf <- mean((preds.rf - test$medv)^2))  
## [1] 4.143477
```

*# The random forest model yields a much lower MSE compared to the linear
regression model.*

- f) Assess the influence of the predictors `rm` and `lstat` in the linear model and the random forest. What do you observe. (R-Hint: `varImpPlot()`, `partialPlot()`)

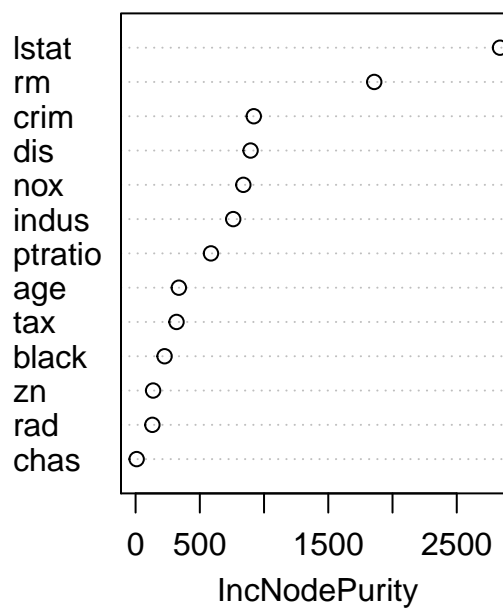
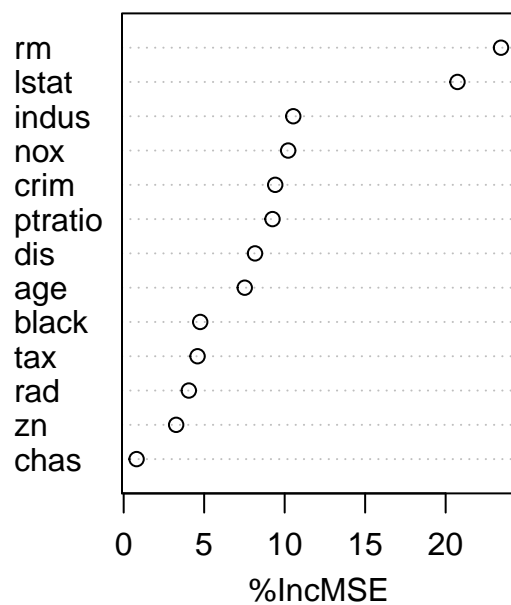
```
summary(fit.lm)
```

```
##  
## Call:  
## lm(formula = medv ~ ., data = train)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max  
## -11.1732  -2.7272  -0.5408   1.6594  23.9450
```



```
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  28.209249   6.094508   4.629 5.24e-06 ***
## crim        -0.070644   0.042379  -1.667 0.096439 .
## zn           0.036046   0.016006   2.252 0.024959 *
## indus        0.030055   0.069986   0.429 0.667874
## chas1        3.419418   0.933343   3.664 0.000288 ***
## nox         -14.582615   4.374454  -3.334 0.000952 ***
## rm           4.856324   0.493995   9.831 < 2e-16 ***
## age         -0.014469   0.015251  -0.949 0.343431
## dis         -1.411125   0.231606  -6.093 3.01e-09 ***
## rad          0.303706   0.075162   4.041 6.59e-05 ***
## tax         -0.013742   0.004219  -3.257 0.001238 **
## ptratio     -0.941851   0.154394  -6.100 2.88e-09 ***
## black        0.010698   0.003084   3.469 0.000590 ***
## lstat       -0.480384   0.061173  -7.853 5.37e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.593 on 340 degrees of freedom
## Multiple R-squared:  0.7778, Adjusted R-squared:  0.7693
## F-statistic: 91.52 on 13 and 340 DF,  p-value: < 2.2e-16
# both variables are highly significant. If rm increases by 1 unit
# (and everything else remains the same), medv increases by 3.84.
# If lstat increases by 1 unit (and everything else remains the same),
# medv decreases by -0.56.
varImpPlot(fit.rf)
```


fit.rf



```
# rm and lstat are the most important variables

# check the marginal dependency of medv on rm and lstat
partialPlot(fit.rf, Boston, x.var = "rm")
```


Partial Dependence on "rm"



If rm increases, medv increases (in a non-linear manner)

```
partialPlot(fit.rf, Boston, x.var = "lstat")
```

Partial Dependence on "lstat"



If lstat increases, medv decreases (in a non-linear manner)

*# This is in concordance with the effect of the variables in the
linear regression model.*